The Big Matrix OLS Jam

*Please email typos / corrections to
`james.h.bisbee@vanderbilt.edu`

Prof. Bisbee

Vanderbilt University

Lecture Date: 2023/11/09

Slides Updated: 2023-12-23

Agenda

1. Matrix Algebra fun!

2. Multiple Regression

3. Controls

Matrix Algebra Fun! Thanks BK!

- Vectors: ordered arrays denoted $\mathbf{v} = (v_1, v_2, \ldots, v_k)$ or

$$\mathbf{v} = \begin{pmatrix} v_1 \\ v_2 \\ \vdots \\ v_k \end{pmatrix}$$

- (Note that some will denote vectors with bold letters, or with $\vec{v}$)
- Addition and subtraction require two vectors of the same length, $\mathbf{u}$ and $\mathbf{v}$, but are then just adding or subtracting the elements

$$\mathbf{u} \pm \mathbf{v} = \begin{pmatrix} u_1 \pm v_1 \\ u_2 \pm v_2 \\ \vdots \\ u_k \pm v_k \end{pmatrix}$$

Vectors

- Multiplication by a constant c is just multiplying each element by c

$$c\mathbf{v} = \begin{pmatrix} cv_1 \\ cv_2 \\ \vdots \\ cv_k \end{pmatrix}$$

- Multiplication of two vectors is called a **dot product**, written $\mathbf{u} \cdot \mathbf{v}$, and translates to multiplying each element in $\mathbf{u}$ by the corresponding element in $\mathbf{v}$ and then adding them all up

$$\begin{aligned} \mathbf{u} \cdot \mathbf{v} &= u_1 v_1 + u_2 v_2 + \cdots + u_k v_k \\ &= \sum_{m=1}^{k} u_m v_m \end{aligned}$$

Matrices

- A matrix is a two-dimensional array with entries in n rows and m columns, called an $n \times m$ matrix

$$\mathbf{A} = \begin{bmatrix} a_{11} & a_{12} & a_{13} \\ a_{21} & a_{22} & a_{23} \\ a_{31} & a_{32} & a_{33} \end{bmatrix}$$

- As with vectors, matrices can be added and subtracted *as long as they are the same dimensions*

$$\mathbf{A} \pm \mathbf{B} = \begin{bmatrix} a_{11} \pm b_{11} & a_{12} \pm b_{12} & a_{13} \pm b_{13} \\ a_{21} \pm b_{21} & a_{22} \pm b_{22} & a_{23} \pm b_{23} \\ a_{31} \pm b_{31} & a_{32} \pm b_{32} & a_{33} \pm b_{33} \end{bmatrix}$$

- As with vectors, matrices multiplied by a constant are straightforward

$$c\mathbf{A} = \begin{bmatrix} ca_{11} & ca_{12} & ca_{13} \\ ca_{21} & ca_{22} & ca_{23} \\ ca_{31} & ca_{32} & ca_{33} \end{bmatrix}$$

Matrices: Transpose

- Transposing: we can "rotate" $n \times m$ matrices into $m \times n$ matrices
  - Meaning that the first row becomes the first column, the second row becomes the second column, etc.
  - Denoted with $\mathbf{A}^\top$ (or sometimes $\mathbf{A}'$)
- For example:

$$\mathbf{A} = \begin{bmatrix} 99 & 73 & 2 \\ 13 & 40 & 41 \end{bmatrix} \qquad \Leftrightarrow \qquad \mathbf{A}^\top = \begin{bmatrix} 99 & 13 \\ 73 & 40 \\ 2 & 41 \end{bmatrix}$$

Matrices: Transpose

- Properties of transposes

$$
\begin{aligned}
(\mathbf{A}^\top)^\top &= \mathbf{A}, \\
(c\mathbf{A})^\top &= c(\mathbf{A}^\top), \\
(\mathbf{A}+\mathbf{B})^\top &= \mathbf{A}^\top + \mathbf{B}^\top, \\
(\mathbf{A}-\mathbf{B})^\top &= \mathbf{A}^\top - \mathbf{B}^\top, \\
(\mathbf{A}\mathbf{B})^\top &= \mathbf{B}^\top \mathbf{A}^\top.
\end{aligned}
$$

- Note that it doesn't make sense to transpose a scalar
    - But also that this means a scalar is always equal to its transpose: $a = a^\top$

Matrix Multiplication

- Refresher: need to multiply an $n \times m$ matrix by an $m \times p$ matrix.
    - **NOTE**: the number of rows in the second matrix must be equal to the number of columns in the first matrix!
- Resulting matrix is an $n \times p$ matrix whose ij'th element is the **dot product** of the i'th row of the first matrix and the j'th column of the second matrix
- Try it: solve $\mathbf{AB}$ where

$$\mathbf{A} = \begin{bmatrix} 2 & 10 \\ 0 & 1 \\ -1 & 5 \end{bmatrix} \qquad \text{and} \qquad \mathbf{B} = \begin{bmatrix} 1 & 4 \\ -1 & 10 \end{bmatrix}$$

$$\mathbf{AB} = \begin{bmatrix} 2 \cdot 1 + 10 \cdot (-1) & 2 \cdot 4 + 10 \cdot 10 \\ 0 \cdot 1 + 1 \cdot (-1) & 0 \cdot 4 + 1 \cdot 10 \\ (-1) \cdot 1 + 5 \cdot (-1) & (-1) \cdot 4 + 5 \cdot 10 \end{bmatrix} = \begin{bmatrix} -8 & 108 \\ -1 & 10 \\ -6 & 46 \end{bmatrix}$$

Matrix Multiplication

- Properties of matrix multiplication
  - **Associative**: $(\mathbf{AB})\mathbf{C} = \mathbf{A}(\mathbf{BC})$
  - **Distributive**: $\mathbf{A}(\mathbf{B} + \mathbf{C}) = \mathbf{AB} + \mathbf{AC}$
  - **NOT** commutative: $\mathbf{AB} \neq \mathbf{BA}$
  - **Transpose Rule**: $(\mathbf{AB})^\top = \mathbf{B}^\top \mathbf{A}^\top$

Matrix Expectations

- Expectations are easily distributed throughout a matrix

$$\mathbf{X} = \begin{bmatrix} x_{11} & x_{12} & x_{13} \\ x_{21} & x_{22} & x_{23} \\ x_{31} & x_{32} & x_{33} \end{bmatrix} \quad ; \quad E(\mathbf{X}) = \begin{bmatrix} E(x_{11}) & E(x_{12}) & E(x_{13}) \\ E(x_{21}) & E(x_{22}) & E(x_{23}) \\ E(x_{31}) & E(x_{32}) & E(x_{33}) \end{bmatrix}$$

Matrix Derivatives

- Consider a matrix equation of the form $\mathbf{y} = \mathbf{A}\mathbf{x}$, meaning that each row is $y_i = a_{1i}x_1 + a_{2i}x_2 + \cdots + a_{ki}x_k$
- In matrix notation:

$$\begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{bmatrix} = \begin{bmatrix} a_{11} & a_{21} & \dots & a_{k1} \\ a_{12} & a_{22} & \dots & a_{k2} \\ \vdots & \vdots & \ddots & \vdots \\ a_{1n} & a_{2n} & \dots & a_{kn} \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_k \end{bmatrix}$$

- To take the partial derivative with respect to $\mathbf{x}$, we go element by element in $\mathbf{y}$: $\frac{\partial y_1}{\partial \mathbf{x}}, \frac{\partial y_2}{\partial \mathbf{x}}, \ldots, \frac{\partial y_n}{\partial \mathbf{x}}$
- But to do THIS, we again go element by element through each value of $\mathbf{x}$, noting that $\frac{\partial y_1}{\partial x_1} = a_{11}$ and $\frac{\partial y_1}{\partial x_2} = a_{21}$, and that $\frac{\partial y_2}{\partial x_1} = a_{12}$ and $\frac{\partial y_2}{\partial x_2} = a_{22}$

Matrix Derivatives

- We can write these in vector form as follows:

$$\frac{\partial y_1}{\partial \mathbf{x}} = \begin{bmatrix} \frac{\partial y_1}{\partial x_1} \\ \frac{\partial y_1}{\partial x_2} \\ \vdots \\ \frac{\partial y_1}{\partial x_k} \end{bmatrix} = \begin{bmatrix} a_{11} \\ a_{21} \\ \vdots \\ a_{k1} \end{bmatrix}; \quad \frac{\partial y_2}{\partial \mathbf{x}} = \begin{bmatrix} \frac{\partial y_2}{\partial x_1} \\ \frac{\partial y_2}{\partial x_2} \\ \vdots \\ \frac{\partial y_2}{\partial x_k} \end{bmatrix} = \begin{bmatrix} a_{12} \\ a_{22} \\ \vdots \\ a_{k2} \end{bmatrix}; \quad \dots \quad \frac{\partial y_n}{\partial \mathbf{x}} = \begin{bmatrix} \frac{\partial y_n}{\partial x_1} \\ \frac{\partial y_n}{\partial x_2} \\ \vdots \\ \frac{\partial y_n}{\partial x_k} \end{bmatrix} = \begin{bmatrix} a_{1n} \\ a_{2n} \\ \vdots \\ a_{kn} \end{bmatrix}$$

- Now let's just combine each of these vectors of derivatives into its own matrix to yield:

$$\frac{\partial \mathbf{y}}{\partial \mathbf{x}} = \begin{bmatrix} a_{11} & a_{12} & \dots & a_{1n} \\ a_{21} & a_{22} & \dots & a_{2n} \\ \vdots & \vdots & \ddots & \vdots \\ a_{k1} & a_{k2} & \dots & a_{kn} \end{bmatrix} = \mathbf{A}^\top$$

Matrix Derivatives

- Thus $\frac{\partial \mathbf{y}}{\partial \mathbf{x}} = \frac{\partial (\mathbf{A}\mathbf{x})}{\partial \mathbf{x}} = \mathbf{A}^\top$
- From this, we can also note that, given $y = \mathbf{a}^\top \mathbf{x}$, $\frac{\partial y}{\partial \mathbf{x}} = \frac{\partial \mathbf{a}^\top \mathbf{x}}{\partial \mathbf{x}} = \mathbf{a}$
- And also, given $y = \mathbf{x}^\top \mathbf{A}\mathbf{x}$, $\frac{\partial y}{\partial \mathbf{x}} = \frac{\partial \mathbf{x}^\top \mathbf{A}\mathbf{x}}{\partial \mathbf{x}} = 2\mathbf{A}\mathbf{x}$
- And finally, given $y = \mathbf{x}^\top \mathbf{A}\mathbf{x}$, $\frac{\partial y}{\partial \mathbf{A}} = \frac{\partial \mathbf{x}^\top \mathbf{A}\mathbf{x}}{\partial \mathbf{A}} = \mathbf{x}\mathbf{x}^\top$

Special Matrices

- **Zero** matrices: $\mathbf{0}$ has all entries as zero
  - NB: $\mathbf{A}_{r\times c} \cdot \mathbf{0}_{c\times n} = \mathbf{0}_{r\times n}$ and $\mathbf{0}_{n\times r} \cdot \mathbf{A}_{r\times c} = \mathbf{0}_{n\times c}$
- **Square** matrices: $n \times n$ size, meaning the same number of rows as columns
- **Symmetric** square matrices: $\mathbf{A} = \mathbf{A}^\top$
- **Diagonal** symmetric square matrices: zeros everywhere except the diagonal: if i are rows and j are columns, $i \neq j$, then $a_{ij} = 0$.
- **Identity** diagonal symmetric square matrices: $\mathbf{I}_n$ is a diagonal matrix where the diagonals are 1s
  - What is

$$\mathbf{A} = \begin{bmatrix} 99 & 73 & 2 \\ 13 & 40 & 41 \end{bmatrix} \quad \cdot \quad \mathbf{I} = \begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix}$$

Matrix Inversion

- In the scalar world, we know we can rewrite a division problem $\frac{a}{b}$ as a multiplication problem $a \times \frac{1}{b} = a \times b^{-1}$
  - b^{-1} is the inverse of b
  - The (obvious) requirement for the inverse is that $b \times b^{-1} = \frac{b}{1} \times \frac{1}{b} = \frac{b}{b} = 1$
- In the matrix world, the inverse of a matrix $\mathbf{A}$ is denoted $\mathbf{A}^{-1}$ and must also satisfy: $\mathbf{A}\mathbf{A}^{-1} = \mathbf{I}_n$
- Some properties!
  - If $\mathbf{C}$ is an inverse of $\mathbf{A}$, then $\mathbf{A}$ is also the inverse of $\mathbf{C}$
  - If $\mathbf{C}$ and $\mathbf{D}$ are both inverses of $\mathbf{A}$, then $\mathbf{C} = \mathbf{D}$
  - The inverse of an inverse of $\mathbf{A}$ is just $\mathbf{A}$: $(\mathbf{A}^{-1})^{-1} = \mathbf{A}$
  - The inverse of $\mathbf{A}^\top$ is the same as the inverse of $\mathbf{A}$, transposed: $(\mathbf{A}^\top)^{-1} = (\mathbf{A}^{-1})^\top$
  - If you have a scalar c multiplied by a matrix $\mathbf{A}$, then $(c\mathbf{A})^{-1} = \frac{1}{c}\mathbf{A}^{-1}$

Matrix Inversion

- To invert a 2×2 matrix, follow this rule:
- For

$$\mathbf{A} = \begin{bmatrix} a & b \\ c & d \end{bmatrix}$$

- Invert using

$$\mathbf{A}^{-1} = \frac{1}{ad - bc} \begin{bmatrix} d & -b \\ -c & a \end{bmatrix}$$

- where $ad - bc$ is known as the **determinant** of the matrix $\mathbf{A}$, so named because it "determines" whether a matrix is invertible.
    - Why would it not be invertible? If $ad - bc = 0$ or $ad = bc$!

Matrix Inversion

- Matrix inversion gets harder with larger matrices...you can learn how to do it manually, but this is where software like `R` comes in handy!

```
A <- matrix(c(2, 1, 3, 4),
            nrow = 2,
            ncol = 2)
A
```

```
##      [,1] [,2]
## [1,]    2    3
## [2,]    1    4
```

- Use the `solve()` function to get the inverse of A

```
A_inv <- solve(A)
A_inv
```

```
##      [,1] [,2]
## [1,]  0.8 -0.6
## [2,] -0.2  0.4
```

Matrix Math in R

- R also can make our lives easier for matrix multiplication...just use %*% instead of the standard *

```r
# Use %*% to do matrix multiplication
A*A_inv # Doesn't work...just does element-by-element multiplication
```

```
##      [,1] [,2]
## [1,]  1.6 -1.8
## [2,] -0.2  1.6
```

```r
A %*% A_inv # Works! We've proved that A_inv is the inverse of A!
```

```
##      [,1] [,2]
## [1,]    1    0
## [2,]    0    1
```

Why all this!?

- It helps us solve systems of equations!
- Back in the day, you probably had lots of practice with these types of things:

$$
\begin{aligned}
2x_1 + x_2 &= 10, \\
2x_1 - x_2 &= -10
\end{aligned}
$$

- You probably learned to solve it various ways (i.e., solve for x_1 first then plug in)
- We can solve with matrix math instead!

$$
\mathbf{A} = \begin{bmatrix} 2 & 1 \\ 2 & -1 \end{bmatrix},
$$

$$
\mathbf{x} = \begin{bmatrix} x_1 \\ x_2 \end{bmatrix},
$$

$$
\mathbf{b} = \begin{bmatrix} 10 \\ -10 \end{bmatrix}
$$

Systems of Equations

- We can re-write the two equations with matrix notation as $\mathbf{A}\mathbf{x} = \mathbf{b}$
- To solve for $\mathbf{x}$, we just invert $\mathbf{A}$ and write $\mathbf{x} = \mathbf{A}^{-1}\mathbf{b}$

```
A <- matrix(c(2, 2, 1, -1),
            nrow = 2,
            ncol = 2)
b <- matrix(c(10,-10),nrow = 2,ncol = 1)

solve(A)%*%b
```

```
##       [,1]
## [1,]     0
## [2,]    10
```

- $x_1 = 0$ and $x_2 = 10$! So easy!
- Note that there is a unique solution for x_1 and x_2 iff $\mathbf{A}$ is invertible
  - If not, there is either no solution or infinitely many solutions

Multiple Regression (Thanks PJE!)

- We can use matrix algebra to help us with **multiple regression** (one outcome with multiple predictors)
  - Note: **multivariate regression** (multiple outcomes) $\neq$ multiple regression
- Let's start with familiar notation and then break it down: $y_i = \beta_0 + \beta_1 x_i + \beta_2 z_i + u_i$
- What does y *look* like? I mean this literally...what is it in a dataset?
  - It is an n-length vector of values $\mathbf{y}$, one for each row in our dataset!
- $\mathbf{x}$ and $\mathbf{z}$ are the same

```
##     respondent_id           y            x            z
## 1               1  1.48840379 -1.270882210 -0.560475647
## 2               2  1.56929669  0.026706220 -0.230177489
## 3               3 -0.51183694  1.312016436  1.558708314
## 4               4  0.19565146 -0.277034208  0.070508391
## 5               5 -1.36595852 -0.822330832  0.129287735
## 6               6 -0.52127462  1.670037262  1.715064987
## 7               7 -1.57350731 -0.323988263  0.460916206
## 8               8 -2.26255920 -2.933003171 -1.265061235
## 9               9  1.27068095 -1.067079372 -0.686852852
```

Multiple Regression

- Let's now look at the data in a different way, from the perspective of a single unit of observation
    - I.e., if we are dealing with a survey of individuals, our data might have some respondent 7 for whom we observe both y_7 as well as x_7 and z_7
- From this perspective, unit 7 is associated with an outcome y_7 (a single value) and then a vector of predictors: $\mathbf{x}_7 = (x_7, z_7)$

```
dat %>% slice(7)
```

```
##   respondent_id         y          x         z
## 1             7 -1.573507 -0.3239883 0.4609162
```

- We can write our regression equation for this specific respondent as $y_7 = \beta_0 + \beta_1 x_7 + \beta_2 z_7 + u_7$, or we can write it as $y_7 = \mathbf{x}_7 \cdot \beta + u_7$
    - β is now itself a **vector** of coefficients: $\beta = (\beta_0, \beta_1, \beta_2)$
    - $\mathbf{x}_7$ now needs to include the number 1: $\mathbf{x}_7 = (1, x_7, z_7)$ in order to capture the β_0 coefficient.

Multiple Regression

- We can then think of β as a $k \times 1$ vector (where k is the number of predictors) and $\mathbf{x}_7$ as a $1 \times k$ vector, and then matrix multiply them!

$$\begin{aligned} y_7 &= \mathbf{x}_7 \cdot \beta + u_7 \\ &= \begin{bmatrix} 1 & x_7 & z_7 \end{bmatrix} \cdot \begin{bmatrix} \beta_0 \\ \beta_1 \\ \beta_2 \end{bmatrix} + u_7 \\ &= \beta_0 + \beta_1 x_7 + \beta_2 z_7 + u_7, \end{aligned}$$

- Now this was just one observation in our data, but we can imagine doing this for every single row, and then stacking our equations on top of each other

$$\begin{aligned} y_1 &= \beta \cdot \mathbf{x}_1 + u_1, \\ y_2 &= \beta \cdot \mathbf{x}_2 + u_2, \\ &\vdots \\ y_n &= \beta \cdot \mathbf{x}_N + u_n. \end{aligned}$$

Multiple Regression

- As with any system of equations, we can re-write as vectors and matrices

$$\mathbf{y} = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{bmatrix}, \qquad \mathbf{X} = \begin{bmatrix} \mathbf{x}_1 \\ \mathbf{x}_2 \\ \vdots \\ \mathbf{x}_n \end{bmatrix} = \begin{bmatrix} 1 & x_1 & z_1 \\ 1 & x_2 & z_2 \\ \vdots & \vdots & \\ 1 & x_n & z_n \end{bmatrix}, \qquad \mathbf{u} = \begin{bmatrix} u_1 \; u_2 \vdots u_n \end{bmatrix}$$

- Plugging in: $\mathbf{y} = \mathbf{X} \cdot \beta + \mathbf{u}$

- Note that this is the same as writing:

$$\begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{bmatrix}_{n \times 1} = \begin{bmatrix} 1 & x_1 & z_1 \\ 1 & x_2 & z_2 \\ \vdots & \vdots & \\ 1 & x_n & z_n \end{bmatrix}_{n \times k} \begin{bmatrix} \beta_0 \\ \beta_1 \\ \beta_2 \end{bmatrix}_{k \times 1} + \begin{bmatrix} u_1 \\ u_2 \\ \vdots \\ u_n \end{bmatrix}_{n \times 1}$$

- where k is the number of parameters (in this case, 3) and n is the number of observations

Multiple Regression

- Note that $\mathbf{y} = \mathbf{X} \cdot \beta + \mathbf{u}$ is assumed to be a reflection of the real world
    - Aside: prove to yourself that $\mathbf{y} = \mathbf{X} \cdot \beta + \mathbf{u}$ and $\mathbf{y} = \beta^\top \cdot \mathbf{X} + \mathbf{u}$ are equivalent
- We estimate these, as before, with our OLS estimators $\hat{\beta}$
- To do so, we first calculate our residuals as $u = y - X\hat{\beta}$, and then add them up and square them.
    - In the **scalar** world, we would write this as $\sum u_i^2$.
    - In the **vector** world, we write this as $\mathbf{u}^\top \mathbf{u}$. Take a moment and try to see why!

$$\mathbf{u}^\top \mathbf{u} = [\, u_1 \quad u_2 \quad \dots \quad u_n \,]_{1\times n} \cdot \begin{bmatrix} u_1 \\ u_2 \\ \vdots \\ u_n \end{bmatrix}_{n\times 1} = [\, u_1 * u_1 + u_2 * u_2 + \cdots + u_n * u_n \,]_{1\times n} = \sum u_i^2$$

Multiple Regression

- We can re-write the sum of squared residuals as $\mathbf{u}^\top \mathbf{u} = (\mathbf{y} - \mathbf{X}\hat{\beta})^\top (\mathbf{y} - \mathbf{X}\hat{\beta})$ by plugging in
- Now let's try doing some reorganizing of this

$$\begin{aligned}(\mathbf{y} - \mathbf{X}\hat{\beta})^\top (\mathbf{y} - \mathbf{X}\hat{\beta}) &= (\mathbf{y}^\top - \hat{\beta}^\top \mathbf{X}^\top)(\mathbf{y} - \mathbf{X}\hat{\beta}) \\ &= \mathbf{y}^\top \mathbf{y} - \mathbf{y}^\top \mathbf{X}\hat{\beta} - \hat{\beta}^\top \mathbf{X}^\top \mathbf{y} + \hat{\beta}^\top \mathbf{X}^\top \mathbf{X}\hat{\beta}\end{aligned}$$

- To subtract, it must be that $\mathbf{y}^\top \mathbf{y}$ is conformable with $\mathbf{y}^\top \mathbf{X}\hat{\beta}$, meaning they must have the same dimensions
- Note that $\mathbf{y}^\top \mathbf{y}$ is a scalar, meaning that $\mathbf{y}^\top \mathbf{X}\hat{\beta}$ must also be a scalar (by conformability)
  - Thus we can re-write $\mathbf{y}^\top \mathbf{X}\hat{\beta} = (\mathbf{y}^\top \mathbf{X}\hat{\beta})^\top = \hat{\beta}^\top \mathbf{X}^\top \mathbf{y}$ (by transpose of a scalar)
- Substitute this in to reduce to:

$$\mathbf{u}^\top \mathbf{u} = \mathbf{y}^\top \mathbf{y} - 2\hat{\beta}^\top \mathbf{X}^\top \mathbf{y} + \hat{\beta}^\top \mathbf{X}^\top \mathbf{X}\hat{\beta}$$

Multiple Regression

- Take the derivative with respect to $\hat{\beta}$ and set it equal to zero, just like we did in the bivariate case

$$
\begin{aligned}
\frac{\partial \mathbf{u}^\top \mathbf{u}}{\partial \hat{\beta}} &= -2\mathbf{X}^\top \mathbf{y} + 2\mathbf{X}^\top \mathbf{X}\hat{\beta} = 0 \\
(\mathbf{X}^\top \mathbf{X})\hat{\beta} &= \mathbf{X}^\top \mathbf{y}
\end{aligned}
$$

- To solve for $\hat{\beta}$, we need to pre-multiply both the left and the right by the inverse of $(\mathbf{X}^\top \mathbf{X})$, assuming it exists

$$
\begin{aligned}
(\mathbf{X}^\top \mathbf{X})\hat{\beta} &= \mathbf{X}^\top \mathbf{y} \\
(\mathbf{X}^\top \mathbf{X})^{-1}(\mathbf{X}^\top \mathbf{X})\hat{\beta} &= (\mathbf{X}^\top \mathbf{X})^{-1}\mathbf{X}^\top \mathbf{y} \\
\mathbf{I}\hat{\beta} &= (\mathbf{X}^\top \mathbf{X})^{-1}\mathbf{X}^\top \mathbf{y} \\
\hat{\beta} &= (\mathbf{X}^\top \mathbf{X})^{-1}\mathbf{X}^\top \mathbf{y}
\end{aligned}
$$

- **Welcome to the matrix definition of the OLS estimator!**

Unbiasedness

- Is this unbiased?
- To start, let's fiddle with the preceding definition of $\hat{\beta}$ a little bit by replacing $\mathbf{y}$ with $\mathbf{X}\beta + \mathbf{u}$.
    - Note that this requires **Assumption 1**: that the population model can be written as $\mathbf{y} = \mathbf{X}\beta + \mathbf{u}$

$$
\begin{aligned}
\hat{\beta} &= (\mathbf{X}^\top\mathbf{X})^{-1}\mathbf{X}^\top\mathbf{y} \\
&= (\mathbf{X}^\top\mathbf{X})^{-1}\mathbf{X}^\top(\mathbf{X}\beta + \mathbf{u}) \\
&= (\mathbf{X}^\top\mathbf{X})^{-1}\mathbf{X}^\top\mathbf{X}\beta + (\mathbf{X}^\top\mathbf{X})^{-1}\mathbf{X}^\top\mathbf{u} \\
&= \mathbf{I}\beta + (\mathbf{X}^\top\mathbf{X})^{-1}\mathbf{X}^\top\mathbf{u} \\
&= \beta + (\mathbf{X}^\top\mathbf{X})^{-1}\mathbf{X}^\top\mathbf{u}
\end{aligned}
$$

Unbiasedness

- Now let's invoke **Assumption 2** that these observations are drawn from an i.i.d. random sample, allowing us take expectations

$$\begin{aligned} E(\hat{\beta}) &= E\left[\beta + (\mathbf{X}^\top \mathbf{X})^{-1}\mathbf{X}^\top \mathbf{u}\right] \\ &= E(\beta) + E\left[(\mathbf{X}^\top \mathbf{X})^{-1}\mathbf{X}^\top \mathbf{u}\right] \\ &= \beta + E\left[(\mathbf{X}^\top \mathbf{X})^{-1}\mathbf{X}^\top \mathbf{u}\right] \end{aligned}$$

- Note that this requires $(\mathbf{X}^\top \mathbf{X})^{-1}$ to exist, so we'll invoke **Assumption 3**: there is no perfect multicollinearity among our X values
  - *Compare this to the non-zero variance assumption invoked when we were working with scalars in the bivariate case*

Unbiasedness

- Finally, let's invoke our most heroic assumption **Assumption 4**: $E(\mathbf{u}|\mathbf{X}) = \mathbf{0}$, and then rely on the law of iterated expectations (LIE)

$$\begin{aligned} E(\hat{\beta} \mid \mathbf{X}) &= \beta + E\Big[(\mathbf{X}^\top\mathbf{X})^{-1}\mathbf{X}^\top\mathbf{u} \mid \mathbf{X}\Big] \\ &= \beta + (\mathbf{X}^\top\mathbf{X})^{-1}\mathbf{X}^\top E(\mathbf{u} \mid \mathbf{X}) \\ &= \beta + (\mathbf{X}^\top\mathbf{X})^{-1}\mathbf{X}^\top \mathbf{0} \\ &= \beta \end{aligned}$$

Properties of the OLS Estimators

- $\mathbf{X}^\top \mathbf{u} = 0$: To prove, substitute the definition of $\mathbf{y} = \mathbf{X}\hat{\beta} + \mathbf{u}$ into the normal equation

$$
\begin{aligned}
(\mathbf{X}^\top \mathbf{X})\hat{\beta} &= \mathbf{X}^\top \mathbf{y} \\
(\mathbf{X}^\top \mathbf{X})\hat{\beta} &= \mathbf{X}^\top (\mathbf{X}\hat{\beta} + \mathbf{u}) \\
(\mathbf{X}^\top \mathbf{X})\hat{\beta} &= (\mathbf{X}^\top \mathbf{X})\hat{\beta} + \mathbf{X}^\top \mathbf{u} \\
0 &= \mathbf{X}^\top \mathbf{u}
\end{aligned}
$$

Properties of the OLS Estimators

- If our regression specification includes a constant, $\sum u_i = 0$: To prove, look inside the matrices!

$$\begin{bmatrix} x_{11} & x_{12} & \dots & x_{1n} \\ x_{21} & x_{22} & \dots & x_{2n} \\ \vdots & \vdots & \ddots & \vdots \\ x_{k1} & x_{k2} & \dots & x_{kn} \end{bmatrix} \cdot \begin{bmatrix} u_1 \\ u_2 \\ \vdots \\ u_n \end{bmatrix} = \begin{bmatrix} x_{11} * u_1 + x_{12} * u_2 + \dots + x_{1n} * u_n \\ x_{21} * u_1 + x_{22} * u_2 + \dots + x_{2n} * u_n \\ \vdots \\ x_{k1} * u_1 + x_{k2} * u_2 + \dots + x_{kn} * u_n \end{bmatrix} = \begin{bmatrix} 0 \\ 0 \\ \vdots \\ 0 \end{bmatrix}$$

- If $\mathbf{X}^\top \mathbf{u} = \mathbf{0}$, then every column $\mathbf{x}_k$'s dot product with $\mathbf{u}$ must be zero
- Since the first column of $\mathbf{X}$ is all 1, then this first column reduces to $\sum u_i = 0$
- Also note that therefore $\bar{u} = 0$ since $\bar{u} = \frac{\sum u_i}{n}$

Properties of the OLS Estimators

- The regression **hyperplane** (no longer a single line, since we have multiple predictors) will pass through $\bar{X}$ and $\bar{y}$
  - We just showed that $\bar{u} = 0$, and we know that $u = y - X\hat{\beta}$
  - Thus $\bar{u} = \bar{y} - \bar{x}\hat{\beta}$, meaning $\bar{y} = \bar{x}\hat{\beta}$
- The predicted values of y are uncorrelated with the residuals
  - $\hat{\mathbf{y}} = \mathbf{X}\hat{\beta}$, meaning that

$$\begin{aligned}\hat{\mathbf{y}}^\top \mathbf{u} &= \mathbf{X}\hat{\beta}^\top \mathbf{u} \\ &= \hat{\beta}^\top \mathbf{X}^\top \mathbf{u} \\ &= \hat{\beta}^\top \cdot \mathbf{0}\end{aligned}$$

Variance in matrix world

- Finally, let's calculate the variance of our OLS estimators, $\hat{\beta}$
- In the scalar world, we calculate the variance of a random variable as $var(x) = E(x - E(x))^2$
- The matrix equivalent of this is called (confusingly) the **covariance** of a random vector, written $cov(\mathbf{x})$
  - Defined as $cov(\mathbf{x}) = E[(\mathbf{x} - E(\mathbf{x}))(\mathbf{x} - E(\mathbf{x}))^\top]$
- Let's write this out!

$$cov(\mathbf{x}) = E\left\{ \begin{bmatrix} x_1 - E(x_1) \\ x_2 - E(x_2) \\ \vdots \\ x_n - E(x_n) \end{bmatrix} \begin{bmatrix} x_1 - E(x_1) & x_2 - E(x_2) & \dots & x_n - E(x_n) \end{bmatrix} \right\}$$

Variance in matrix world

$$cov(\mathbf{x}) = E\left\{ \begin{bmatrix} (x_1 - E(x_1))^2 & (x_1 - E(x_1))(x_2 - E(x_2)) & \dots & (x_1 - E(x_1))(x_n - E(x_n)) \\ (x_2 - E(x_2))(x_1 - E(x_1)) & (x_2 - E(x_2))^2 & \dots & (x_2 - E(x_2))(x_n - E(x_n)) \\ \vdots & \vdots & \ddots & \vdots \\ (x_n - E(x_n))(x_1 - E(x_1)) & (x_n - E(x_n))(x_2 - E(x_2)) & \dots & (x_n - E(x_n))^2 \end{bmatrix} \right\}$$

- Distribute expectations throughout to get

$$cov(\mathbf{x}) = \begin{bmatrix} \sigma^2_{x_1} & cov(x_1, x_2) & \dots & cov(x_1, x_n) \\ cov(x_2, x_1) & \sigma^2_{x_2} & \dots & cov(x_2, x_n) \\ \vdots & \vdots & \ddots & \vdots \\ cov(x_n, x_1) & cov(x_n, x_2) & \dots & \sigma^2_{x_n} \end{bmatrix}$$

- NB: this is called the covariance matrix of the random vector $\mathbf{x}$, AKA the **variance-covariance** matrix
  - Often depicted with $\mathbf{\Sigma}$

Variance of $\hat{\beta}$

- So now let's use this to calculate the **variance of** $\hat{\beta}$
  - Note that we have already demonstrated that $E(\hat{\beta}) = \beta$
- Also note that $\hat{\beta} = \beta + (\mathbf{X}^\top\mathbf{X})^{-1}\mathbf{X}^\top\mathbf{u}$, or $\hat{\beta} - \beta = (\mathbf{X}^\top\mathbf{X})^{-1}\mathbf{X}^\top\mathbf{u}$
- Plug in

$$
\begin{aligned}
E[(\hat{\beta} - \beta)(\hat{\beta} - \beta)^\top] &= E\Bigg[\bigg((\mathbf{X}^\top\mathbf{X})^{-1}\mathbf{X}^\top\mathbf{u}\bigg)\bigg((\mathbf{X}^\top\mathbf{X})^{-1}\mathbf{X}^\top\mathbf{u}\bigg)^\top\Bigg] \\
&= E\Bigg[\bigg((\mathbf{X}^\top\mathbf{X})^{-1}\mathbf{X}^\top\mathbf{u}\bigg)\bigg(\mathbf{u}^\top\mathbf{X}(\mathbf{X}^\top\mathbf{X})^{-1}\bigg)\Bigg] \\
&= E\Bigg[(\mathbf{X}^\top\mathbf{X})^{-1}\mathbf{X}^\top\mathbf{u}\mathbf{u}^\top\mathbf{X}(\mathbf{X}^\top\mathbf{X})^{-1}\Bigg]
\end{aligned}
$$

Errors

- This is the variance of our estimator: $E\left[(\mathbf{X}^\top\mathbf{X})^{-1}\mathbf{X}^\top\mathbf{u}\mathbf{u}^\top\mathbf{X}(\mathbf{X}^\top\mathbf{X})^{-1}\right]$
- Taking a step back:
  - We have a **L**inear **E**stimator $\hat{\beta}$
  - We have proved it is **U**nbiased
- Is it the "best"? (Remember, **B**est **L**inear **U**nbiased **E**stimator is **BLUE**)
- To prove it is BLUE, we require **Assumption 5**: $E(\mathbf{u}\mathbf{u}^\top \mid \mathbf{X}) = \sigma^2\mathbf{I}$. AKA: "spherical errors"
  - a. **Homoskedasticity**: $var(u_1) = var(u_2) = \cdots = var(u_n) = \sigma^2$
  - b. **No autocorrelation**: $cov(u_i, u_j) = 0$ for all $i \neq j$

Errors

- Let's write out *Assumption 5**:

$$E(\mathbf{u}\mathbf{u}^\top \mid \mathbf{X}) = E\left(\begin{bmatrix} u_1 \mid \mathbf{X} \\ u_2 \mid \mathbf{X} \\ \vdots \\ u_n \mid \mathbf{X} \end{bmatrix} \begin{bmatrix} u_1 \mid \mathbf{X} & u_2 \mid \mathbf{X} & \dots & u_n \mid \mathbf{X} \end{bmatrix} \right)$$

$$= E \begin{bmatrix} u_1^2 \mid \mathbf{X} & u_1 u_2 \mid \mathbf{X} & \dots & u_1 u_n \mid \mathbf{X} \\ u_2 u_1 \mid \mathbf{X} & u_2^2 \mid \mathbf{X} & \dots & u_2 u_n \mid \mathbf{X} \\ \vdots & \vdots & \ddots & \vdots \\ u_n u_1 \mid \mathbf{X} & u_n u_2 \mid \mathbf{X} & \dots & u_n^2 \mid \mathbf{X} \end{bmatrix}$$

Errors

- Distribute expectations to get:

$$E(\mathbf{u}\mathbf{u}^\top \mid \mathbf{X}) = \begin{bmatrix} E(u_1^2 \mid \mathbf{X}) & E(u_1 u_2 \mid \mathbf{X}) & \dots & E(u_1 u_n \mid \mathbf{X}) \\ E(u_2 u_1 \mid \mathbf{X}) & E(u_2^2 \mid \mathbf{X}) & \dots & E(u_2 u_n \mid \mathbf{X}) \\ \vdots & \vdots & \ddots & \vdots \\ E(u_n u_1 \mid \mathbf{X}) & E(u_n u_2 \mid \mathbf{X}) & \dots & E(u_n^2 \mid \mathbf{X}) \end{bmatrix}$$

- From **Assumption 5**:
  - Homoskedasticity states that the variance of $u_i = \sigma^2$ for all i, or $VAR(u_i|\mathbf{X}) = \sigma^2 \ \ \forall\, i$
  - No autocorrelation states that $cov(u_i, u_j|\mathbf{X}) = 0$

Errors

- Thus, assumption 5 allows us to re-write:

$$
\begin{aligned}
E(\mathbf{u}\mathbf{u}^\top \mid \mathbf{X}) &= \begin{bmatrix} E(u_1^2 \mid \mathbf{X}) & E(u_1 u_2 \mid \mathbf{X}) & \dots & E(u_1 u_n \mid \mathbf{X}) \\ E(u_2 u_1 \mid \mathbf{X}) & E(u_2^2 \mid \mathbf{X}) & \dots & E(u_2 u_n \mid \mathbf{X}) \\ \vdots & \vdots & \ddots & \vdots \\ E(u_n u_1 \mid \mathbf{X}) & E(u_n u_2 \mid \mathbf{X}) & \dots & E(u_n^2 \mid \mathbf{X}) \end{bmatrix} \\
&= \begin{bmatrix} \sigma^2 & 0 & \dots & 0 \\ 0 & \sigma^2 & \dots & 0 \\ \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \dots & \sigma^2 \end{bmatrix}
\end{aligned}
$$

- which is the same as writing $\sigma^2\mathbf{I}$

Variance of $\hat{\beta}$

- So we have $E\left[(\mathbf{X}^\top\mathbf{X})^{-1}\mathbf{X}^\top\mathbf{u}\mathbf{u}^\top\mathbf{X}(\mathbf{X}^\top\mathbf{X})^{-1}\right]$
- Take the LIE conditional on $\mathbf{X}$ to get

$$\begin{aligned}
E[(\hat{\beta}-\beta)(\hat{\beta}-\beta)^\top \mid \mathbf{X}] &= E\left[(\mathbf{X}^\top\mathbf{X})^{-1}\mathbf{X}^\top\mathbf{u}\mathbf{u}^\top\mathbf{X}(\mathbf{X}^\top\mathbf{X})^{-1} \,\middle|\, \mathbf{X}\right] \\
&= (\mathbf{X}^\top\mathbf{X})^{-1}\mathbf{X}^\top E(\mathbf{u}\mathbf{u}^\top \mid \mathbf{X})\mathbf{X}(\mathbf{X}^\top\mathbf{X})^{-1} \\
&= (\mathbf{X}^\top\mathbf{X})^{-1}\mathbf{X}^\top \;(\sigma^2\mathbf{I})\; \mathbf{X}(\mathbf{X}^\top\mathbf{X})^{-1} \\
&= \sigma^2\mathbf{I}(\mathbf{X}^\top\mathbf{X})^{-1}\mathbf{X}^\top\mathbf{X}(\mathbf{X}^\top\mathbf{X})^{-1} \\
&= \sigma^2\mathbf{I}\mathbf{I}(\mathbf{X}^\top\mathbf{X})^{-1} \\
&= \sigma^2(\mathbf{X}^\top\mathbf{X})^{-1}
\end{aligned}$$

What does this give us?

- The OLS estimator -- $\hat{\beta}$ -- is a random vector, distributed with mean β and a variance-covariance matrix $\sigma^2(\mathbf{X}^\top\mathbf{X})^{-1}$
- We might be particularly interested in just one of the coefficients contained within this vector (i.e., β_1 speaks to a theoretical quantity of interest, while the other $\beta_2, \beta_3, \ldots, \beta_k$ are controls)
- To find the mean of $\hat{\beta}_1$, we look inside our vector of expected values of $\hat{\beta}$ and extract the element corresponding to $E(\hat{\beta}_1) = (\mathbf{X}^\top\mathbf{X}_1^{-1}\mathbf{X}_1^\top\mathbf{y})$
- To find the variance of $\hat{\beta}_1$, we look inside our variance-covariance matrix $cov(\hat{\beta}) = \mathbf{\Sigma}_{\hat{\beta}}$ and extract the entry corresponding to $E(\hat{\beta}_1 - E(\hat{\beta}_1))^2 = \sigma^2(\mathbf{X}^\top\mathbf{X})^{-1}_{11}$
- As always, we never know σ^2, meaning we never really know $var(\hat{\beta})$
- In practice, we estimate the unknown σ^2 with $\hat{\sigma}^2 = \frac{\mathbf{u}^\top\mathbf{u}}{n-k}$
  - Note that we are assuming k includes β_0. If not, we write as $\hat{\sigma}^2 = \frac{\mathbf{u}^\top\mathbf{u}}{n-k-1}$

A few final comments

- As in the univariate and bivariate cases, we can appeal to the **C**entral **L**imit **T**heorem (CLT) to assume that the sampling distribution of $\hat{\beta} \xrightarrow{d} MVN(\beta, \sigma^2(\mathbf{X}^\top\mathbf{X})^{-1})$
    - The symbol $\xrightarrow{d}$ means "distributed asymptotically as"
- The multivariate normal (MVN) joint distribution means that we can extract any element of $\hat{\beta}$ and standardize it, and it will be distributed asymptotically as the standard normal
    - I.e., $\frac{\hat{\beta}_k - \bar{\beta}_k}{\sqrt{\hat{\sigma}^2(\mathbf{X}^\top\mathbf{X})^{-1}_{kk}}} \xrightarrow{d} \mathcal{N}(0,1)$
    - NB: the statistic $\frac{\hat{\beta}_k - \bar{\beta}_k}{\sqrt{\hat{\sigma}^2(\mathbf{X}^\top\mathbf{X})^{-1}_{kk}}}$ is distributed according to the Student's t distribution with $N-K-1$ degrees of freedom: $\frac{\hat{\beta}_k - \bar{\beta}_k}{\sqrt{\hat{\sigma}^2(\mathbf{X}^\top\mathbf{X})^{-1}_{kk}}} \sim t_{N-k-1}$
- In small samples, we make **one more assumption** that the errors are normally distributed

FWL and Partialling Out

- To understand what multiple regression looks like in matrix form, we need some helper concepts
- The "residual maker" is a matrix $\mathbf{M}$ that, when multiplied by $\mathbf{y}$, creates **residuals $\mathbf{u}$**
- Start with the definition of the residual: $\mathbf{u} = \mathbf{y} - \hat{\mathbf{y}}$ and substitute $\hat{\mathbf{y}} = \mathbf{X}\hat{\beta}$ in

$$\mathbf{u} = \mathbf{y} - \mathbf{X}\hat{\beta}$$

- Now replace $\hat{\beta}$ with the definition $(\mathbf{X}^\top\mathbf{X})^{-1}\mathbf{X}^\top\mathbf{y}$

$$\begin{aligned}\mathbf{u} &= \mathbf{y} - \mathbf{X}(\mathbf{X}^\top\mathbf{X})^{-1}\mathbf{X}^\top\mathbf{y} \\ &= (\mathbf{I} - \mathbf{X}(\mathbf{X}^\top\mathbf{X})^{-1}\mathbf{X}^\top)\mathbf{y} \\ &= \mathbf{M}\mathbf{y}\end{aligned}$$

FWL and Partialling Out

- $\mathbf{M}$ is super helpful. It is both square and **idempotent**, meaning that $\mathbf{MM} = \mathbf{M}$. (Try proving this for yourself!)
- It also has the properties:
  1. $\mathbf{MX} = \mathbf{0}$
  2. $\mathbf{Mu} = \mathbf{u}$

FWL and Partialling Out

- The "hat" matrix is a matrix $\mathbf{H}$ that, when multiplied by $\mathbf{y}$, creates **predicted values** $\hat{\mathbf{y}}$

$$\begin{aligned}\hat{\mathbf{y}} &= \mathbf{y} - \mathbf{u} \\ &= (\mathbf{I} - \mathbf{M})\mathbf{y} \\ &= \mathbf{H}\mathbf{y}\end{aligned}$$

- So we now have $\mathbf{M} = \mathbf{I} - \mathbf{X}(\mathbf{X}^\top\mathbf{X})^{-1}\mathbf{X}^\top$ and $\mathbf{H} = \mathbf{I} - \mathbf{M}$
- But this just means $\mathbf{H} = \mathbf{X}(\mathbf{X}^\top\mathbf{X})^{-1}\mathbf{X}^\top$

FWL and Partialling Out

- These "hat" and "residual maker" matrices can help us understand OVB and, more generally, what "controlling" for a variable means in the matrix world
- Consider the classic example where the true regression equation is given by $\mathbf{y} = \mathbf{x}_1\beta_1 + \mathbf{x}_2\beta_2 + u$, but we mistakenly omit $\mathbf{x}_2$
- The true normal equation is:

$$\begin{bmatrix} \mathbf{x}_1^\top \mathbf{x}_1 & \mathbf{x}_1^\top \mathbf{x}_2 \\ \mathbf{x}_2^\top \mathbf{x}_1 & \mathbf{x}_2^\top \mathbf{x}_2 \end{bmatrix}^{-1} \cdot \begin{bmatrix} \mathbf{x}_1\mathbf{y} \\ \mathbf{x}_2\mathbf{y} \end{bmatrix} = \begin{bmatrix} \hat{\beta}_1 \\ \hat{\beta}_2 \end{bmatrix}$$

FWL and Partialling Out

- First, solve for $\hat{\beta}_1$

$$\begin{aligned}(\mathbf{x}_1^\top \mathbf{x}_1)\hat{\beta}_1 + (\mathbf{x}_1^\top \mathbf{x}_2)\hat{\beta}_2 &= \mathbf{x}_1^\top \mathbf{y} \\ (\mathbf{x}_1^\top \mathbf{x}_1)\hat{\beta}_1 &= \mathbf{x}_1^\top \mathbf{y} - (\mathbf{x}_1^\top \mathbf{x}_2)\hat{\beta}_2 \\ \hat{\beta}_1 &= (\mathbf{x}_1^\top \mathbf{x}_1)^{-1}\mathbf{x}_1^\top \mathbf{y} - (\mathbf{x}_1^\top \mathbf{x}_1)^{-1}(\mathbf{x}_1^\top \mathbf{x}_2)\hat{\beta}_2 \\ \hat{\beta}_1 &= (\mathbf{x}_1^\top \mathbf{x}_1)^{-1}\mathbf{x}_1^\top (\mathbf{y} - \mathbf{x}_2\hat{\beta}_2)\end{aligned}$$

- Recognize this?
    - $(\mathbf{x}_1^\top \mathbf{x}_1)^{-1}\mathbf{x}_1^\top \mathbf{x}_2$ is the regression of $\mathbf{x}_2$ on $\mathbf{x}_1$. This will be zero if $\mathbf{x}_2$ is unrelated to $\mathbf{x}_1$
    - $\hat{\beta}_2$ is the relationship between $\mathbf{y}$ and $\mathbf{x}_2$.
- **This is just OVB in matrix form**

FWL and Partialling Out

- Now let's see what happens when we control for $\mathbf{x}_2$
- Start with $\hat{\beta}_1 = (\mathbf{x}_1^\top \mathbf{x}_1)^{-1} \mathbf{x}_1^\top (\mathbf{y} - \mathbf{x}_2 \hat{\beta}_2)$
- Then do direct multiplication on the second row in

$$\begin{bmatrix} \mathbf{x}_1^\top \mathbf{x}_1 & \mathbf{x}_1^\top \mathbf{x}_2 \\ \mathbf{x}_2^\top \mathbf{x}_1 & \mathbf{x}_2^\top \mathbf{x}_2 \end{bmatrix}^{-1} \cdot \begin{bmatrix} \mathbf{x}_1 \mathbf{y} \\ \mathbf{x}_2 \mathbf{y} \end{bmatrix} = \begin{bmatrix} \hat{\beta}_1 \\ \hat{\beta}_2 \end{bmatrix}$$

- to yield $\mathbf{x}_2^\top \mathbf{x}_1 \hat{\beta}_1 + \mathbf{x}_2^\top \mathbf{x}_2 \hat{\beta}_2 = \mathbf{x}_2^\top \mathbf{y}$
- Finally, substitute in our definition of $\hat{\beta}_1$ to get $\mathbf{x}_2^\top \mathbf{x}_1 (\mathbf{x}_1^\top \mathbf{x}_1)^{-1} \mathbf{x}_1^\top \mathbf{y} - \mathbf{x}_2^\top \mathbf{x}_1 (\mathbf{x}_1^\top \mathbf{x}_1)^{-1} \mathbf{x}_1^\top \mathbf{x}_2 \hat{\beta}_2 + \mathbf{x}_2^\top \mathbf{x}_2 \hat{\beta}_2 = \mathbf{x}_2^\top \mathbf{y}$

FWL and Partialling Out

- So this is horrible, but try this!

$$
\begin{aligned}
\mathbf{x}_2^\top \mathbf{y} - \mathbf{x}_2^\top \mathbf{x}_1(\mathbf{x}_1^\top \mathbf{x}_1)^{-1}\mathbf{x}_1^\top \mathbf{y} &= \mathbf{x}_2^\top \mathbf{x}_2 \hat{\beta}_2 - \mathbf{x}_2^\top \mathbf{x}_1(\mathbf{x}_1^\top \mathbf{x}_1)^{-1}\mathbf{x}_1^\top \mathbf{x}_2 \hat{\beta}_2 \\
\mathbf{x}_2^\top \mathbf{y} - \mathbf{x}_2^\top \mathbf{x}_1(\mathbf{x}_1^\top \mathbf{x}_1)^{-1}\mathbf{x}_1^\top \mathbf{y} &= [\mathbf{x}_2^\top \mathbf{x}_2 - \mathbf{x}_2^\top \mathbf{x}_1(\mathbf{x}_1^\top \mathbf{x}_1)^{-1}\mathbf{x}_1^\top \mathbf{x}_2]\hat{\beta}_2 \\
\mathbf{x}_2^\top \mathbf{y} - \mathbf{x}_2^\top \mathbf{x}_1(\mathbf{x}_1^\top \mathbf{x}_1)^{-1}\mathbf{x}_1^\top \mathbf{y} &= [(\mathbf{x}_2^\top - \mathbf{x}_2^\top \mathbf{x}_1(\mathbf{x}_1^\top \mathbf{x}_1)^{-1}\mathbf{x}_1^\top)\mathbf{x}_2]\hat{\beta}_2 \\
\mathbf{x}_2^\top \mathbf{y} - \mathbf{x}_2^\top \mathbf{x}_1(\mathbf{x}_1^\top \mathbf{x}_1)^{-1}\mathbf{x}_1^\top \mathbf{y} &= [(\mathbf{x}_2^\top (\mathbf{I} - \mathbf{x}_1(\mathbf{x}_1^\top \mathbf{x}_1)^{-1}\mathbf{x}_1^\top)\mathbf{x}_2]\hat{\beta}_2 \\
(\mathbf{x}_2^\top - \mathbf{x}_2^\top \mathbf{x}_1(\mathbf{x}_1^\top \mathbf{x}_1)^{-1}\mathbf{x}_1^\top)\mathbf{y} &= [(\mathbf{x}_2^\top (\mathbf{I} - \mathbf{x}_1(\mathbf{x}_1^\top \mathbf{x}_1)^{-1}\mathbf{x}_1^\top)\mathbf{x}_2]\hat{\beta}_2 \\
\mathbf{x}_2^\top (\mathbf{I} - \mathbf{x}_1(\mathbf{x}_1^\top \mathbf{x}_1)^{-1}\mathbf{x}_1^\top)\mathbf{y} &= [(\mathbf{x}_2^\top (\mathbf{I} - \mathbf{x}_1(\mathbf{x}_1^\top \mathbf{x}_1)^{-1}\mathbf{x}_1^\top)\mathbf{x}_2]\hat{\beta}_2 \\
\hat{\beta}_2 &= [(\mathbf{x}_2^\top (\mathbf{I} - \mathbf{x}_1(\mathbf{x}_1^\top \mathbf{x}_1)^{-1}\mathbf{x}_1^\top)\mathbf{x}_2]^{-1}\mathbf{x}_2^\top (\mathbf{I} - \mathbf{x}_1(\mathbf{x}_1^\top \mathbf{x}_1)^{-1}\mathbf{x}_1^\top)\mathbf{y} \\
&= (\mathbf{x}_2^\top \mathbf{M}_1 \mathbf{x}_2)^{-1}(\mathbf{x}_2^\top \mathbf{M}_1 \mathbf{y})
\end{aligned}
$$

FWL and Partialling Out

- So we now have $\hat{\beta}_2 = (\mathbf{x}_2^\top \mathbf{M}_1 \mathbf{x}_2)^{-1}(\mathbf{x}_2^\top \mathbf{M}_1 \mathbf{y})$
- Remember that $\mathbf{M}$ is the residual maker, meaning that $\mathbf{M}_1$ is making residuals for regressions on the $\mathbf{x}_1$ variables
  - $\mathbf{M}_1\mathbf{y}$ therefore creates residuals from regressing $\mathbf{y}$ on $\mathbf{x}_1$
  - $\mathbf{M}_1\mathbf{x}_2$ therefore creates residuals from regressing $\mathbf{x}_2$ on $\mathbf{x}_1$
- Since $\mathbf{M}$ is both idempotent and symmetric, we can rewrite as $\hat{\beta}_2 = (\mathbf{x}_2^{*\top}\mathbf{x}_2)^{-1}\mathbf{x}_2^{*\top}\mathbf{y}^*$
  - Where $\mathbf{x}_2^* = \mathbf{M}_1\mathbf{x}_2$ and $\mathbf{y}^* = \mathbf{M}_1\mathbf{y}$
- This leads to the **Frisch-Waugh-Lovell** Theorem: In the OLS regression of a vector $\mathbf{y}$ on two sets of variables $\mathbf{x}_1$ and $\mathbf{x}_2$, $\hat{\beta}_2$ is the coefficient obtained when the residuals from a regression of $\mathbf{y}$ on $\mathbf{x}_1$ alone are regressed on the set of residuals obtained when $\mathbf{x}_2$ is regressed on $\mathbf{x}_1$

FWL and Partialling Out

- Imagine the following model (reverting back to the layperson's notation here):
  $Y = \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \beta_3 X_3 + u$

- According to FWL:

  1. Regress Y on X_1 and obtain the residuals $\hat{u}_1$ (i.e., $\mathbf{M}_1\mathbf{y}$ in matrix notation)

  2. Regress X_2 on X_1 and obtain the residuals $\hat{u}_2$ (i.e., $\mathbf{M}_1\mathbf{x}_2$ in matrix notation)

  3. Regress X_3 on X_1 and obtain the residuals $\hat{u}_3$ (i.e., $\mathbf{M}_1\mathbf{x}_3$ in matrix notation)

  4. Regress $\hat{u}_1$ on $\hat{u}_2$ and $\hat{u}_3$: $\hat{u}_1 = \rho_0 + \rho_1\hat{u}_2 + \rho_2\hat{u}_3 + \epsilon$

- $\hat{\beta}_2$ will be equal to $\hat{\rho}_1$ and $\hat{\beta}_3$ will be equal to $\hat{\rho}_3$!

- Steps 2 and 3 are called "partialling out" or "netting out" the effect of X_1. For this reason, the coefficients in multiple regression are often referred to as "partial regression coefficients".

FWL and Partialling Out

- Let's try it!

```
X1 <- rnorm(100)
X2 <- rnorm(100)
X3 <- rnorm(100)

# True beta_1 = 1, beta_2 = -1, beta_3 = 3
Y <- X1 - X2 + 3*X3 + rnorm(100)

# Multiple regression
mFull <- lm(Y ~ X1 + X2 + X3)

# FWL way
u_1 <- resid(lm(Y ~ X1))
u_2 <- resid(lm(X2 ~ X1))
u_3 <- resid(lm(X3 ~ X1))
mRes <- lm(u_1 ~ u_2 + u_3)
```

FWL and Partialling Out

- As promised, we get the same estimates for $\hat{\beta}_2$ and $\hat{\beta}_3$ whether we estimate them in the standard multiple regression setting, or if we use the FWL residualizer approach

```
# Same coefficients!
round(coef(mFull)[c(3,4)],4)
```

```
##      X2      X3 
## -0.9926  2.9206
```

```
round(coef(mRes)[c(2,3)],4)
```

```
##     u_2     u_3 
## -0.9926  2.9206
```